

Bias-Variance Tradeoff and Regularization

Sourav Karmakar


Outline

1 Bias-Variance Tradeoff

2 Regularization

Introduction: The Goal of a Model

In machine learning, we train a model on a **training dataset** to predict an output $\hat{y}$ for a given input x .



Our ultimate goal is **generalization**: the model should not only perform well on the training data but also on unseen **test data**. Two main sources of error prevent our models from generalizing perfectly: **Bias** and **Variance**.

What is Bias?

Bias is the error introduced by approximating a complex real-world problem by a simpler model. Bias represents the error stemming from overly simplistic assumptions in a learning algorithm. It measures how far off our predicted values are from the true values, *on average*.

Mathematical Formulation:

$$\text{Bias}(\text{Model}) = \mathbb{E}[\hat{y}] - y$$

Where:

- $\hat{y}$ is the model's prediction.
- $\mathbb{E}[\hat{y}]$ is the *expected* (or average) prediction of the model if we trained it over and over again on many different samples of training data.
- y is the true, actual output.

Numerical Example: Bias

Let's calculate the bias of a model predicting house prices.

Scenario: The true price of a specific house is **\$500k** ($y = 500$). We train our model on 3 different random samples of data and predict the price for this house.

Numerical Example: Bias

Let's calculate the bias of a model predicting house prices.

Scenario: The true price of a specific house is **\$500k** ($y = 500$). We train our model on 3 different random samples of data and predict the price for this house.

Step 1: Get predictions ($\hat{y}$)

- Dataset 1 Prediction: 300
- Dataset 2 Prediction: 320
- Dataset 3 Prediction: 280

Numerical Example: Bias

Let's calculate the bias of a model predicting house prices.

Scenario: The true price of a specific house is **\$500k** ($y = 500$). We train our model on 3 different random samples of data and predict the price for this house.

Step 1: Get predictions ($\hat{y}$)

- Dataset 1 Prediction: 300
- Dataset 2 Prediction: 320
- Dataset 3 Prediction: 280

Step 2: Expected prediction $\mathbb{E}[\hat{y}]$

$$\mathbb{E}[\hat{y}] = \frac{300 + 320 + 280}{3} = \mathbf{300}$$

Numerical Example: Bias

Let's calculate the bias of a model predicting house prices.

Scenario: The true price of a specific house is **\$500k** ($y = 500$). We train our model on 3 different random samples of data and predict the price for this house.

Step 1: Get predictions ($\hat{y}$)

- Dataset 1 Prediction: 300
- Dataset 2 Prediction: 320
- Dataset 3 Prediction: 280

Step 2: Expected prediction $\mathbb{E}[\hat{y}]$

$$\mathbb{E}[\hat{y}] = \frac{300 + 320 + 280}{3} = \mathbf{300}$$

Step 3: Calculate Bias

$$\text{Bias} = 300 - 500 = \mathbf{-200}$$

Numerical Example: Bias

Let's calculate the bias of a model predicting house prices.

Scenario: The true price of a specific house is **\$500k** ($y = 500$). We train our model on 3 different random samples of data and predict the price for this house.

Step 1: Get predictions ($\hat{y}$)

- Dataset 1 Prediction: 300
- Dataset 2 Prediction: 320
- Dataset 3 Prediction: 280

Step 2: Expected prediction $\mathbb{E}[\hat{y}]$

$$\mathbb{E}[\hat{y}] = \frac{300 + 320 + 280}{3} = \mathbf{300}$$

Step 3: Calculate Bias

$$\text{Bias} = 300 - 500 = \mathbf{-200}$$

Conclusion: This model has a **High Bias**. It consistently underestimates the true value by a large margin.

What is Variance?

Variance is the amount that the model's prediction changes if we train it on a different training dataset. It measures the *inconsistency* or variability of the predictions.

Mathematical Formulation:

$$\text{Variance} = \mathbb{E} [(\hat{y} - \mathbb{E}[\hat{y}])^2]$$

We want to reduce variance because we want our model to be **robust**—it should not drastically change its predictions just because the training data changed slightly.

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Model B (A new, complex model):

Predictions: 100, 900, 500

Mean = 500 (*Low Bias*: $500 - 500 = 0$)

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Model B (A new, complex model):

Predictions: 100, 900, 500

Mean = 500 (*Low Bias*: $500 - 500 = 0$)

$$\text{Var}_B = \frac{(100-500)^2 + (900-500)^2 + (500-500)^2}{3}$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Model B (A new, complex model):

Predictions: 100, 900, 500

Mean = 500 (*Low Bias*: $500 - 500 = 0$)

$$\text{Var}_B = \frac{(100-500)^2 + (900-500)^2 + (500-500)^2}{3}$$

$$\text{Var}_B = \frac{(-400)^2 + (400)^2 + 0}{3}$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Model B (A new, complex model):

Predictions: 100, 900, 500

Mean = 500 (*Low Bias*: $500 - 500 = 0$)

$$\text{Var}_B = \frac{(100-500)^2 + (900-500)^2 + (500-500)^2}{3}$$

$$\text{Var}_B = \frac{(-400)^2 + (400)^2 + 0}{3}$$

$$\text{Var}_B \approx \mathbf{106,666} \quad (\text{High Variance})$$

Numerical Example: Variance

Let's calculate the variance for two different models predicting the same \$500k house.

Model A (From previous slide):

Predictions: 300, 320, 280

Mean ($\mathbb{E}[\hat{y}]$) = 300

$$\text{Var}_A = \frac{(300-300)^2 + (320-300)^2 + (280-300)^2}{3}$$

$$\text{Var}_A = \frac{0+400+400}{3}$$

$$\text{Var}_A = \mathbf{266.67} \quad (\text{Low Variance})$$

Model B (A new, complex model):

Predictions: 100, 900, 500

Mean = 500 (*Low Bias*: $500 - 500 = 0$)

$$\text{Var}_B = \frac{(100-500)^2 + (900-500)^2 + (500-500)^2}{3}$$


$$\text{Var}_B = \frac{(-400)^2 + (400)^2 + 0}{3}$$

$$\text{Var}_B \approx \mathbf{106,666} \quad (\text{High Variance})$$


Conclusion: Model B has zero bias but **High Variance**. Its predictions are wildly unstable depending on the training data.

Visualizing Bias and Variance (The Dart Board Diagram)


Imagine the center of the target (bullseye) in the dart board is the true value y . The dots are predictions from models trained on different training samples.




Low Bias
Low Variance
(Ideal)



High Bias
Low Variance



Low Bias
High Variance



High Bias
High Variance

The Total Error Equation

The expected Mean Squared Error (MSE) of a model on unseen data can be mathematically decomposed into three parts:

$$\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$

- **Bias²:** Error due to wrong assumptions in the model.
- **Variance:** Error due to the model's sensitivity to small fluctuations in the training set.
- **Irreducible Error (ϵ):** The noise inherent in the problem itself. No model can ever reduce this error.

High Bias Causes and Effects

What causes High Bias?

- Using overly **simpler models** (e.g., fitting a straight line to curved data).
- Having too few features.

Effects on Error:

- **Training Error:** HIGH (The model cannot even capture the underlying pattern in the training data).
- **Test Error:** HIGH.

This situation is famously known as **Underfitting**. The model is too rigid to learn anything useful.

High Variance Causes and Effects

What causes High Variance?

- Using very **complex models** (e.g., a high-degree polynomial curve).
- Having too many features compared to the number of data points.


Effects on Error:

- **Training Error:** LOW (The model perfectly memorizes the training data, including the random noise).
- **Test Error:** HIGH (Because it memorized noise, it fails completely on new data).

This situation is famously known as **Overfitting**.

The Bias-Variance Tradeoff Curve


We cannot minimize both Bias and Variance simultaneously. Decreasing bias usually increases variance, and vice versa.




Goal: Find the "sweet spot" (optimum complexity) that minimizes the Total Error on the test dataset.

Underfitting vs. Overfitting Revisited


Let's visualize the tradeoff using a dataset.



Underfitting (High Bias)



Optimal Fit



Overfitting (High Variance)

How do we stop an overfitted model from drawing those extreme curves? We use **Regularization**.

The Role of Model Parameters (θ)

In parametric models like Linear or Logistic Regression, the output $\hat{y}$ is determined by **weights** (parameters) θ :

$$\hat{y} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_k x_k$$

Interpretation of θ_j :

- $\theta_j = \frac{\Delta \hat{y}}{\Delta x_j}$. It represents the change in response per unit change in feature x_j .
- If a weight θ_j is **very large**, the model's response changes drastically for even tiny fluctuations (noise) in x_j .
- **Large weights = High Sensitivity = High Variance.**

What is Regularization?

Regularization is a technique to reduce overfitting by limiting (shrinking) the parameters (θ). We force the model to keep the weights small.

We do this by introducing a “**Penalty**” term into the standard Loss function:

$$L(\Theta, \alpha) = \text{Loss}(\Theta) + \alpha \cdot (\text{Penalty})$$

- **Loss(Θ):** Can be MSE (Regression), BCE, or CCE (Classification).
- **Penalty:** A function of the model weights ($f(\theta_1, \theta_2 \dots)$).
- **α (Regularization Parameter):** Controls how much penalty is applied.
 - If $\alpha = 0$: No penalty (Unregularized).
 - If α is high: Heavy penalty (Weights approach 0).

5. L2 Regularization and Ridge Regression

L2 Regularization is a general technique that adds the squared magnitude of the weights (L2-norm) as a penalty to *any* loss function.

Ridge Regression is the specific model created when we apply L2 Regularization to standard Linear Regression.

$$L(\Theta, \alpha) = \underbrace{\frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2}_{\text{Standard MSE Loss}} + \underbrace{\frac{\alpha}{2} \sum_{j=1}^k \theta_j^2}_{\text{L2 Penalty Term}}$$

Why don't we regularize the bias term (θ_0)?

5. L2 Regularization and Ridge Regression

L2 Regularization is a general technique that adds the squared magnitude of the weights (L2-norm) as a penalty to *any* loss function.

Ridge Regression is the specific model created when we apply L2 Regularization to standard Linear Regression.

$$L(\Theta, \alpha) = \underbrace{\frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2}_{\text{Standard MSE Loss}} + \underbrace{\frac{\alpha}{2} \sum_{j=1}^k \theta_j^2}_{\text{L2 Penalty Term}}$$

Why don't we regularize the bias term (θ_0)?

- The bias term θ_0 simply shifts the function up or down.
- It does not determine the "curvature" or the model's sensitivity to input features x .
- Penalizing θ_0 doesn't reduce variance; it just forces the curve away from the data's true baseline.

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Step 2: Update Rule $(\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \eta \frac{\partial L}{\partial \theta_j})$

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Step 2: Update Rule $(\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \eta \frac{\partial L}{\partial \theta_j})$

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \left[\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} + \alpha \theta_j^{(t)} \right]$$

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Step 2: Update Rule $(\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \eta \frac{\partial L}{\partial \theta_j})$

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \left[\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} + \alpha \theta_j^{(t)} \right]$$

Step 3: Rearrange to see the magic!

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Step 2: Update Rule $(\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \eta \frac{\partial L}{\partial \theta_j})$

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \left[\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} + \alpha \theta_j^{(t)} \right]$$

Step 3: Rearrange to see the magic!

$$\theta_j^{(t+1)} = (1 - \eta\alpha)\theta_j^{(t)} - \eta [\text{Standard Gradient}]$$

Gradient Descent for Ridge Regression

To train the model, we use Gradient Descent. We need the derivative of the new loss function with respect to θ_j .

Step 1: Calculate the Gradient

$$\frac{\partial L}{\partial \theta_j} = \left(\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} \right) + \alpha \theta_j$$

Step 2: Update Rule $(\theta_j^{(t+1)} \leftarrow \theta_j^{(t)} - \eta \frac{\partial L}{\partial \theta_j})$

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \left[\frac{1}{m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)}) x_j^{(i)} + \alpha \theta_j^{(t)} \right]$$


Step 3: Rearrange to see the magic!

$$\theta_j^{(t+1)} = (1 - \eta\alpha) \theta_j^{(t)} - \eta [\text{Standard Gradient}]$$

*Notice $(1 - \eta\alpha)$: Before applying the standard update, we shrink the current weight by a small fraction. This is why L2 is also called **Weight Decay**.*

Visualizing Loss with Contour Diagrams


What is a Contour Diagram? It is a 2D representation of a 3D bowl-like surface. Each ring (contour) represents points where the Loss value is the same. The center is the minimum loss.



**Unregularized MSE
Loss Contours**

L2 Regularization Contours

In Ridge Regression, we want to minimize the MSE, but we are **constrained** by the L2 penalty budget: $\theta_1^2 + \theta_2^2 \leq t$. This constraint forms a **Circle**.



The solution is where the MSE ellipse first touches the L2 circle. Weights are shrunk towards zero, but rarely become exactly zero.

L1 Regularization and LASSO

L1 Regularization is a general technique that adds the absolute sum of the weights (L1-norm) as a penalty to *any* loss function.

LASSO (**L**east **A**bsolute **S**hrinkage and **S**election **O**perator) is the specific model created when L1 Regularization is applied to standard Linear Regression.

$$L(\Theta, \alpha) = \underbrace{\frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2}_{\text{Standard MSE Loss}} + \underbrace{\alpha \sum_{j=1}^k |\theta_j|}_{\text{L1 Penalty Term}}$$


Gradient Descent for LASSO:

Because $|\theta_j|$ is not differentiable at 0, the update rule uses the **sign** of θ_j :

$$\theta_j^{(t+1)} = \theta_j^{(t)} - \eta \left[\text{Standard Gradient} + \alpha \cdot \text{sign}(\theta_j^{(t)}) \right]$$

L1 Regularization Contours and Feature Selection

The L1 constraint ($|\theta_1| + |\theta_2| \leq t$) forms a **Diamond** shape with sharp corners on the axes.



How LASSO Performs Feature Selection

Why does LASSO zero out features?

- Look at the L1 constraint: It has sharp corners that lie **exactly on the axes** (e.g., where $\theta_2 = 0$).
- The expanding elliptical contours of the MSE loss are highly likely to hit these sharp corners first.
- When the intersection happens on an axis, the corresponding weight becomes **exactly zero**.
- By forcing unimportant feature weights to exactly zero, LASSO automatically performs **Feature Selection**, creating sparse, interpretable models.

7. Combined L1 & L2 and Elastic-Net

What if we want the feature selection of L1, but the stability of L2? We can combine both penalties into a general loss function.


Elastic-Net is the specific model created when this combined penalty is applied to standard Linear Regression.

$$L(\Theta, \alpha, \beta) = \underbrace{\frac{1}{2m} \sum_{i=1}^m (\hat{y}^{(i)} - y^{(i)})^2}_{\text{Standard MSE Loss}} + \underbrace{\alpha \left[\beta \sum_{j=1}^k |\theta_j| + \frac{1 - \beta}{2} \sum_{j=1}^k \theta_j^2 \right]}_{\text{Combined Penalty Term}}$$

- α : Overall regularization strength.
- β : The **L1 ratio** parameter ($0 \leq \beta \leq 1$).
 - If $\beta = 0$: Only L2 is applied (Ridge).
 - If $\beta = 1$: Only L1 is applied (Lasso).

Elastic-Net Contours

The constraint shape of Elastic-Net is a compromise between the sharp diamond of L1 and the smooth circle of L2. It looks like a "bulgy" diamond.



It retains some sharp corners for feature selection, but the curve provides grouping effects (correlated features are selected or dropped together).

Summary

- **Bias** (Underfitting): Model is too simple. Train and test errors are both high.
- **Variance** (Overfitting): Model is too complex. Train error is low, but test error is high.
- **Tradeoff**: We must find the optimal complexity to minimize total error on unseen data.
- **Regularization**: Adds a penalty term to the loss function to shrink weights and reduce Variance.
- **L2 Regularization & Ridge**: L2 adds a squared penalty (circular contours). Ridge applies this to Linear Regression for smooth weight shrinkage.
- **L1 Regularization & LASSO**: L1 adds an absolute penalty (diamond contours). LASSO applies this to Linear Regression, automatically performing **feature selection**.
- **Combined Penalty & Elastic-Net**: Balances L1 and L2 constraints to achieve both feature selection and model stability.

Thank You